

Real-Time Linux



Agenda

- **Real Time Linux**
- **Real Time Applications Interface (RTAI)**
 - Hardware Abstraction Layer
 - Real Time Schedulers
 - Real Time Services
 - Features and Performance
- **Programming in RTAI**
 - Application Development

Real Time Systems - Recap

- **A real-time system is a “System capable of guaranteeing **logical correctness & timing requirements** of the processes under its control”**
 - Fast – Low latency (quick response to external, asynchronous events and context switching)
 - Predictable – deterministic completion time of tasks with certainty

Real-Time Systems - Recap

- **Real Time Operating Systems are typically used for control or communications applications in which untimely response can have catastrophic consequences**
 - Air traffic controllers
 - Missile guidance systems
 - Health monitoring systems
- **Scheduling guarantees are offered by real-time operating systems such as QNX, LynxOS or VxWorks**


Standard Linux OS

- Linux is a **time-sharing OS**
 - Good performance as a **desktop/server OS** with Sophisticated services
 - Designed to **maximize throughput** and give a fair share of the CPU
 - Scheduler **avoids process starvation**
 - **User space preemptible**
 - Not preemptible when executing system calls(in Kernel space)
 - **Non deterministic** timing behaviour of some kernel services: memory allocation, system calls
 - **Changes in Interrupt latency**
 - Disabling Interrupts : spinlocks, evil drivers masking out interrupts
 - Shared interrupts : all handlers are executed (don't use shared IRQ lines for critical events)

Real time Linux extensions

- **Goal of Real time extensions**

- Add **extra layer between hardware** and the Linux kernel to manage real-time tasks separately
- **Changes in the kernel sources** which ensures real time platform, with low latency and high predictability requirements
- Strict separation of hard and non hard real time processes
- **Communication and synchronization** services to allow an effective interaction between the two environments



Real-Time Linux - The Concept

- In a **real-time Linux** system
 - Real-time kernel has a **higher priority than the Linux** kernel
 - Real-time tasks are executed by the real-time kernel, and normal Linux programs are allowed to run when no real-time tasks have to be executed
 - **Linux is treated as the idle task** of the real-time scheduler
 - When this idle task runs, it executes its own scheduler and schedules the normal Linux processes.
 - A normal Linux process is preempted when a real-time task becomes ready to run

Real-Time Linux - The Concept

- Scheduler execution
 - Driven by timer interrupts of a clock to reschedule at certain times
 - An executing program can block or voluntary give up the CPU in which case the scheduler is informed by means of a software interrupt
 - Driven by hardware generated interrupts, interrupt the normal scheduled work
- RT Linux uses the flow of interrupts to give the real-time kernel a higher priority
 - When an interrupt arrives, it is first given to the real-time kernel. Interrupts are stored, for Linux to handle, when the real-time kernel is done
 - Real-time kernel can run its real-time tasks driven by these interrupts
- Each flavor of RT Linux does this in its own way

Variants of Real-Time Linux Extensions

- **Real Time Linux (RTLinux)**
 - Developed at the New Mexico Institute of Mining and Technology
 - RTOS Micro kernel running entire Linux in fully preemptive mode
 - Runs special real-time tasks and interrupt handlers
 - FiFo, Shared memory, Semaphores. POSIX mutexes and threads
 - Avg latency - 15us
 - Used to control robots, data acquisition systems, manufacturing plants, and other time-sensitive instruments and machines
- **Xenomai Framework**
 - Platforms - x86, ARM, POWER, IA-64, Blackfin, nios
 - Implementing and migrating real time applications, based on standard APIs or emulators of proprietary RTOS interfaces, such as VxWorks and pSOS
 - Linux-hosted dual kernel, with pure Adeos patch
 - User space and kernel space RT tasks
 - Avg Interrupt response time - 43us

Real Time Application Interface (RTAI)

Real-Time Application Interface (RTAI)

- RTAI Core system

- RTAI is integrated into Linux through a **kernel patch** and a series of add on programs (**loadable modules**) expanding the Linux kernel to hard real time

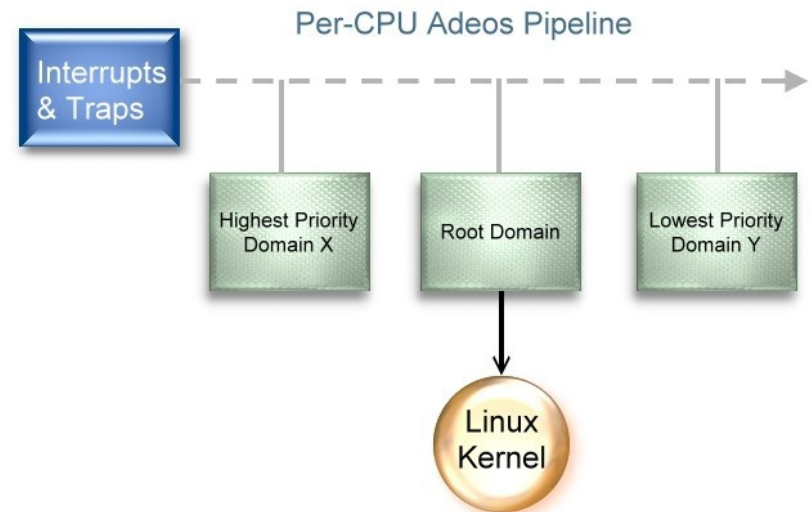
- Adaptive Domain Environment for operating Systems (ADEOS) patch

- **Nanokernel Hardware Abstraction Layer (HAL)**
- Provides flexible **environment for sharing hardware resources** among multiple operating systems
- Implements a **pipeline scheme** into which every domain (OS) has an entry with a predefined priority.
- For each event (interrupts, exceptions, syscalls, ...), the various **domains may handle the event or pass it down the pipeline**
- **RTAI is the highest priority domain** which always processes interrupts before the Linux domain, thus serving any hard real time activity fully preempting anything that is not hard real time

Real-Time Application Interface (RTAI)

- Interrupt pipeline(I-pipe)

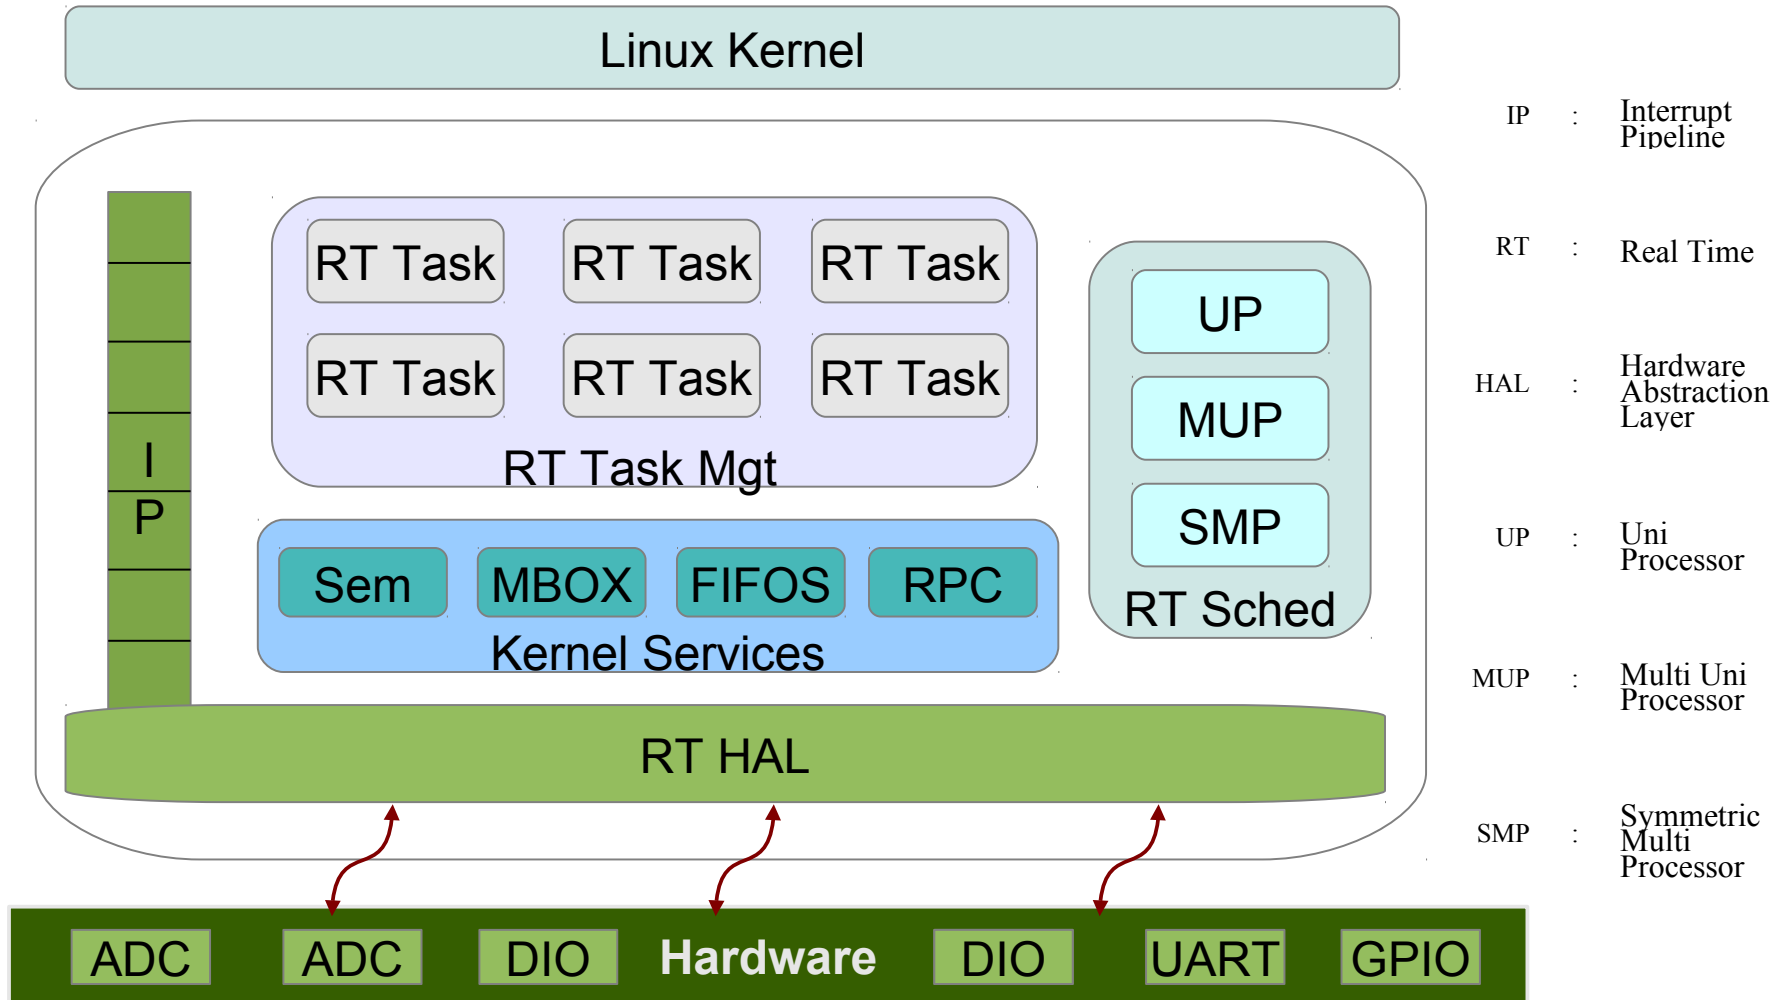
- The high priority domain is at the beginning of the pipeline, so events are delivered first to it
- There is a pipeline for each CPU
- The Linux domain is always the root domain. Other domains are started by the root domain
- Linux starts and loads the kernel modules that implement other domains



Real-Time Application Interface (RTAI)

- Virtualized interrupts disabling
 - Each domain may be “**stalled**”, meaning that it does not accept interrupts
 - **Hardware interrupts are not disabled**(except for the leading domain in pipeline, i.e RTAI) however, instead the interrupts received during that time are logged and replayed when the domain is unstalled
 - Thus **Linux is not permitted to disable hardware interrupts**, and hence, cannot add latency to the interrupt response time of the real time system

RTAI Architecture



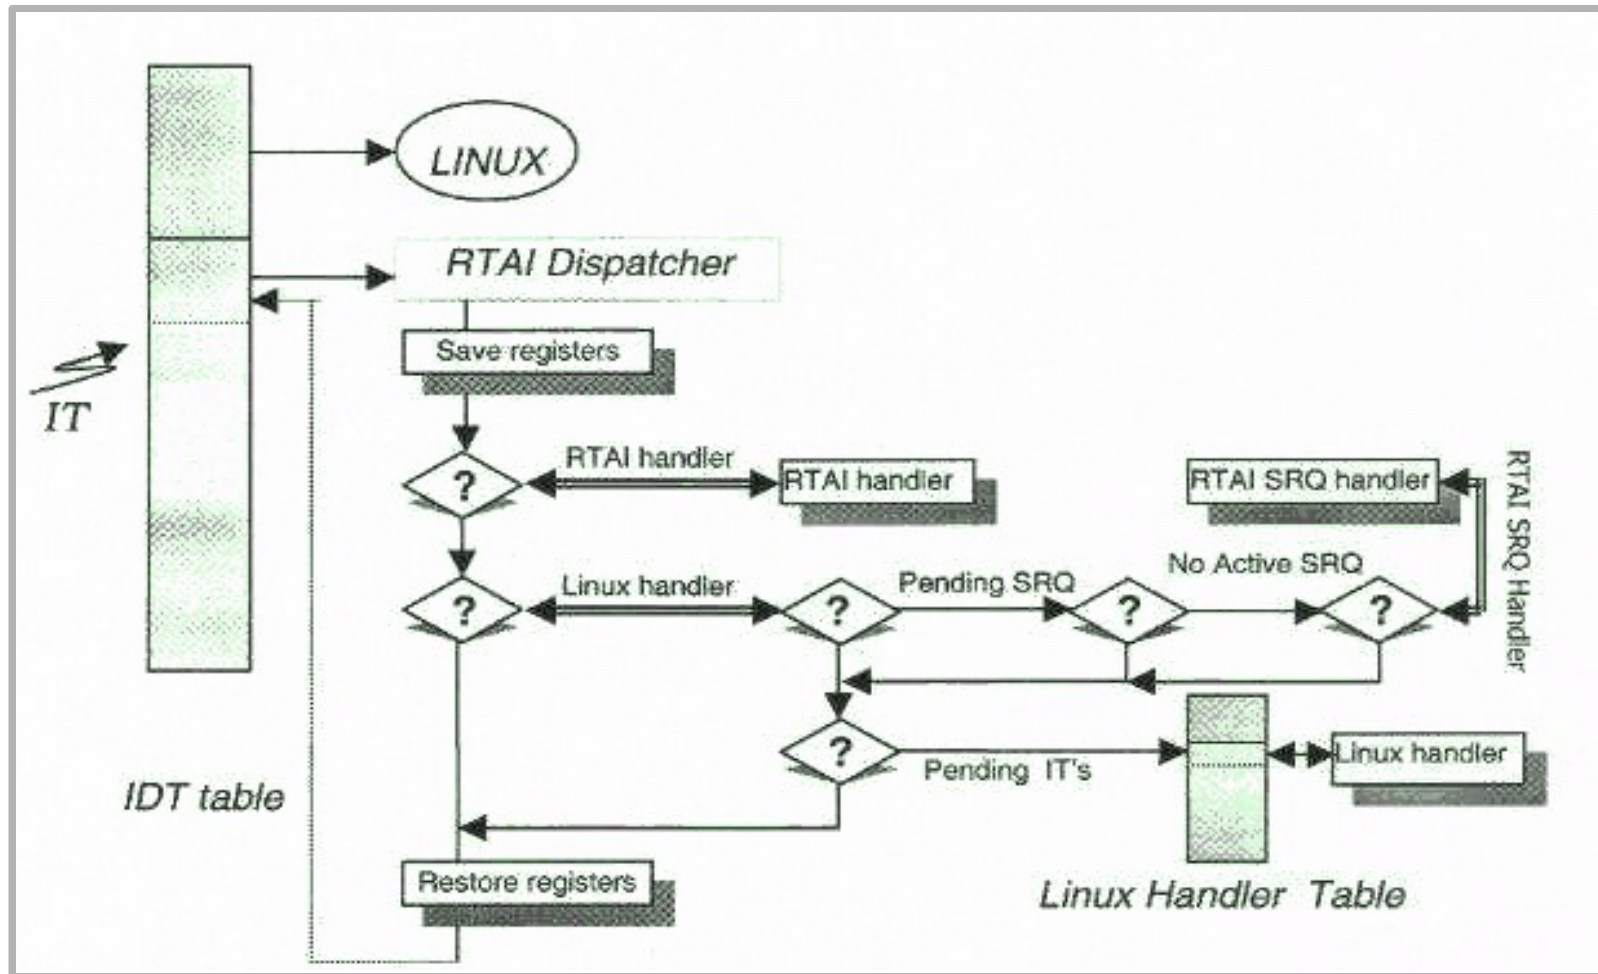
RTHAL

- RTHAL contains
 - Pointer to **Interrupt Descriptor Table (IDT)**
 - A key component which provides a set of pointers, which defines to which processes each of the interrupts should be routed
 - Functions to **enable/disable x86 CPU interrupts** (cli, sti)
 - Functions to mask/unmask interrupt controller
- Upon activation of RTHAL
 - A **new IDT**, which is a duplicate copy of the standard Linux interrupt descriptor table becomes the valid table
 - Handler functions in the new IDT table are changed to the **RTAI interrupt dispatcher** function so that it takes control of the system hardware and its interrupts

RTAI Interrupt Dispatcher

- The **Interrupt Dispatcher** called by the IDT when interrupts occur
 - **Activates the right handler** depending on the owner of the driver, either RTAI / Linux / both (RTAI first then Linux)
 - The addresses of the **Linux handlers are saved** before changing the IDT in the HAL
 - Takes into account the “**SRQ(service request) handlers**”
 - SRQ is a handler used for a system call implementation. SRQ can be either RTAI handler (e.g.: fifo) or User handler
 - RTAI manages up to **31 SRQ's**. SRQ handler addresses are stored in IDT

RTAI Interrupt Dispatcher



RTAI Scheduler

- The RTAI multitasking scheduler, uses [Interrupt-driven](#), Priority-based task scheduling
- Scheduler elects the highest priority task in READY state
 - Priority 0 is the highest priority
 - **0x3fffFfff** is the lowest for rtai tasks
 - Linux is given priority **0x7fffFfff**
- **Pure Periodic Mode (Default)**
 - Program the timer only once with a frequency
 - Task activations are only possible at multiples of the timer period
- **One-shot Mode**
 - Program the timer once for every iteration
 - Arbitrary timings for task activations

RTAI Scheduler Timing

- RTAI schedulers, timed by three type of hard timers
 - 8254 Timer IC
 - Allows one shot and periodic modes
 - Special care is taken in maintaining Linux timing
 - One-Shot mode is provided by Time Stamp Clock (TSC) or Counter2
 - APIC
 - Allows one shot and periodic modes
 - One-Shot mode is provided Time Stamp Clock (TSC)

RTAI Schedulers

- **The RTAI distribution includes three different priority based, pre-emptive real time schedulers**
 - Uni-Processor (UP) scheduler
 - Multi Uni-Processor (MUP) scheduler
 - Symmetric Multi-Processor (SMP) scheduler
- **During the installation process, a scheduler is determined based on the hardware configuration of the target machine**
 - It is then copied and linked so that it is called by the generic `rtai_sched.ko` reference

RTAI Schedulers

- **UP scheduler - For uni-processor platforms**
 - Based on the 8254 based timer
 - Supports either one-shot or periodic scheduling but not both simultaneously
- **SMP scheduler - For multi-processor machines**
 - Based on either 8254 timer or APIC timers
 - Supports either one-shot or periodic scheduling but not both simultaneously
 - Tasks can run symmetrically on any or a cluster of CPUs, or be bound to a single CPU
 - By default all tasks are defined to run on any of the CPUs and are automatically moved between CPUs as the system's processing and load requirements change.

RTAI Schedulers

- **Multi-Uniprocessor scheduler - For multiprocessor platforms only**
 - Supports both one-shot and periodic scheduling simultaneously.
 - The main advantage is the ability to be able to use mixed timers simultaneously, i.e. periodic and one-shot timers.
 - Like the SMP schedulers, the MUP can use inter-CPU services related to semaphores, messages and mailboxes.

RTAI Schedulers

- **All RTAI schedulers incorporate standard RTOS scheduling services like**
 - Resume
 - Yield
 - Suspend
 - Make Periodic
 - Wait Until

RTAI Task States

- The scheduler looks for ready tasks or tasks with expired period.
- A task may be in the following states
 - READY
 - SUSPENDED
 - DELAYED
 - SEMAPHORE
 - SEND
 - RECEIVE
 - RPC
 - RETURN
 - RUNNING

RTAI Features

- Features

- Supports UniProcessor, Multi-UniProcessor and Symmetric Multi- Processor (SMP)
- One-shot and Periodic schedulers
- IPC : Semaphores, mailboxes, FIFOs, shared memory, and RPCs
- POSIX 1003.1c (Pthreads, mutexes and condition variables) & POSIX 1003.1b (Pqueues only) compatibility
- LXRT – which allows the use of the RTAI system calls from within standard user space
- /proc interface – which provides information on the real-time tasks, modules, services and processes extending the standard Linux /proc file system support

Real-Time Application Interface (RTAI)

- The Real Time kernel, all its component parts, and the real time applications are all run in Linux kernel address space as kernel modules
 - These modules can be removed from the kernel on completion of the real time system operation. Advantage of this is, it aids system modularity
 - If the scheduler is unsuitable for a particular application, then the scheduler module can be replaced by one that meets the needs of the application
 - One of the main disadvantages of running in Linux kernel address space is that a bug in a real time task can crash the whole system

Real-Time Application Interface (RTAI)

- Provides
 - Deterministic and pre-emptive performance
 - Allows the use of all standard Linux drivers, applications and functions
 - Highly Modularised
- Uses Linux
 - System and device initialisation
 - Blocking dynamic resource allocation
 - Loadable module mech: to install components of the Real-Time kernel
- Decoupled from the Linux kernel
 - Real-time kernel can be kept small and simple
 - Kernels can be optimized independently

RTAI Performance

- RTAI's performance is very competitive with the commercial Real Time Operating Systems such as VxWorks, QNX etc
- RTAI Offers
 - Context switch times of 4 uSec
 - Interrupt Response of 20 uSec
 - Periodic Tasks with periodicity 100 KHz

RTAI Programming

Topics

- Starting the Timer and Real-Time Scheduler
- Task Initialization
- Scheduling One-shot and Periodic Tasks
- Task Function
- Task APIs

Starting the Timer and Realtime Scheduler

- The timer can be started in either
 - Pure Periodic Mode (Default)
 - Program the timer only once with a frequency
 - Task activations are only possible at multiples of the timer period
 - **rt_set_periodic_mode()**
 - One-shot Mode
 - Program the timer once for every iteration
 - Arbitrary timings for task activations
 - **rt_set_one-shot_mode()**

Note: Mode must be set before using any time related function

Starting the Timer and Realtime Scheduler

- Start the real-time timer and scheduler using
 - **RTIME start_rt_timer(RTIME period)**
 - period is required only for the periodic mode
 - Given in internal count units (clock ticks)
 - RTIME nano2counts(int nanoseconds)
 - Scheduler is activated in the timer interrupt handler
- Stop the timer and scheduler
 - **void stop_rt_timer(void)**


Task Initialization

- **RT_TASK** data structure holds the information about a task
 - Task function, any initial argument passed to it
 - Size of the stack allocated for its variables
 - Task priority
 - Whether or not task uses floating-point math
 - "signal handler" that will be called when the task becomes active

Task Initialization

- **Initializing the Task Structure**(RT_TASK)
 - **rt_task_init (RT_TASK *task, void *rt_thread, int data, int stack_size, int priority, int uses_fp, void *sig_handler)**
 - The newly created real time task is initially in a **suspended state**
- To **start the task** immediately or change a task state from suspended to ready
 - **int rt_task_resume(RT_TASK *task)**

Scheduling Tasks

- Start after some delay, as a periodic task 
 - **int rt_task_make_periodic(RT_TASK *task, RTIME start_time, RTIME period);**
 - start_time - absolute time to begin execution in clock ticks. Typically "now"(i.e., calling rt_get_time())
 - Period - task's period in clock ticks
 - **int rt_task_make_periodic_relative_ns (RT_TASK *task, RTIME start_delay, RTIME period);**
 - start_delay - Delay after which task starts execution, relative to the current time in nanoseconds
 - period - dummy value for a oneshot task, task's period for a periodic task

Task Function

- Periodic task function
 - Should enter an **infinite loop**, in which it does its job
 - On completion of task's activity **wait for its next scheduled cycle**
 - **void rt_task_wait_period(void);**

- Periodic task function code snippet

```
void task_function(int arg) {  
    while (1) {  
        /* Do your thing here */  
        rt_task_wait_period();  
    }  
    return;  
}
```

Task APIs

- Delete a task
 - **int rt_task_delete (RT_TASK *task)**
 - One-shot tasks delete themselves
- Yield current task
 - **void rt_task_yield (void)**
 - Stops and puts the task at the end of the ready list
- Suspend a task
 - **int rt_task_suspend (RT_TASK *task)**
 - Task will not execute until it is resumed again

THANK U