

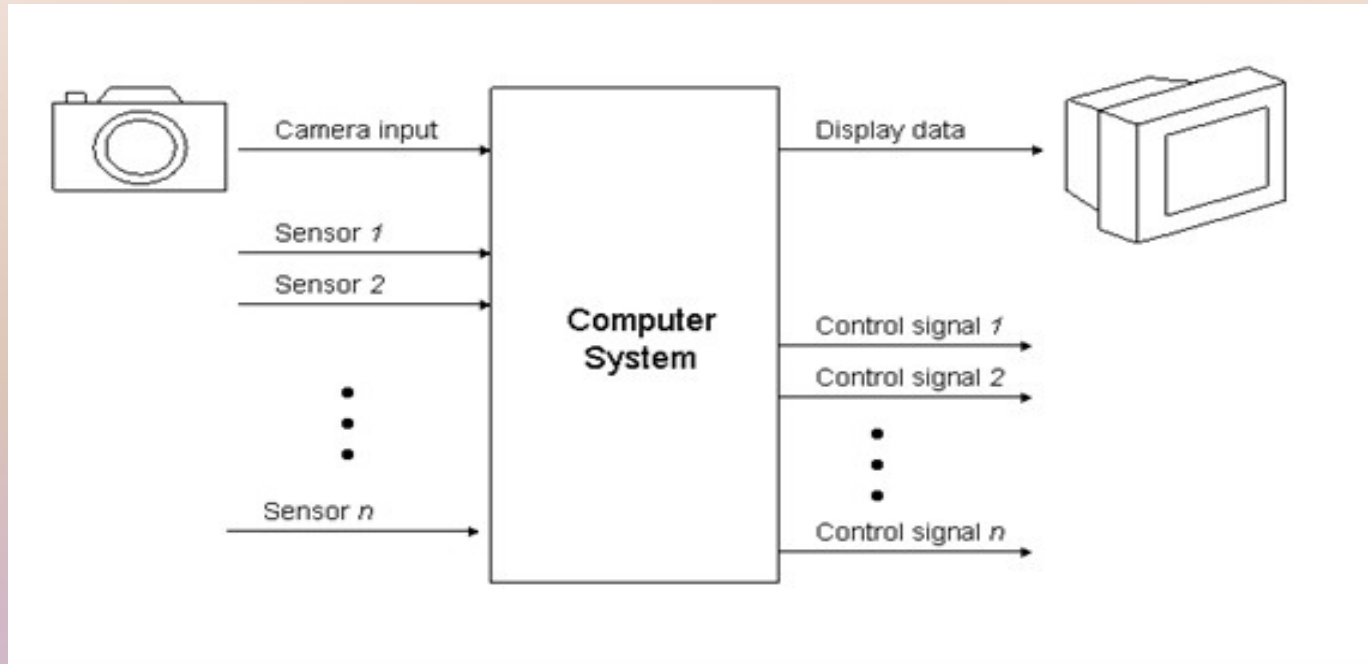
RTOS Concepts & RTAI

Sowjanya
C-DAC, Hyderabad

Agenda

- Real Time Systems
- RTOS Vs GPOS
- RTOS Services
- Real Time Linux Variants
- Real Time Application Interface(RTAI)

Real Time Systems



- Typical real-time control system including inputs from sensors and imaging devices and producing control signals and display information.

Real Time Systems

- **Correctness**

- The correctness of the system depends not only on the logical result, but also on the time at which the results are produced.

- **Predictability**

- Possible to show at “design time” that all the timing constraints of the application will be met.

- **Deterministic**

- For each possible state and each set of inputs, a unique set of outputs and next state of the system can be determined.

Real Time Systems

- Classification of RT systems
 - **Soft real-time system**
 - Performance is degraded but not destroyed by failure to meet response-time constraints.
 - Eg: Multimedia, Interactive video games
 - **Firm real-time system**
 - Missing more than few deadlines, may lead to catastrophe
 - Eg: Robot weed killer
 - **Hard real-time system**
 - Failure to meet a single deadline may lead to catastrophic failure
 - Eg: Aircraft Control Systems, Nuclear Power Stations, Chemical Plants, Life support systems

Real Time Tasks

- **Periodic tasks**

- Time-driven, characteristics are known a priori (p_i , c_i)
- Eg: Task monitoring temperature of a patient in an ICU

- **Aperiodic tasks**

- Event-driven, characteristics not known a priori (a_i , r_i , c_i , d_i)
- Task activated on detecting change in patient's condition

- **Sporadic Tasks**

- Aperiodic tasks with known minimum inter-arrival time.

p_i : task period a_i : arrival time r_i : ready time d_i : deadline c_i : worst case exec time

Real Time Tasks

- The CPU utilization or time-loading factor, is a measure of the percentage of non-idle processing
 - Utilization factor ui for a Task ti is (ci/pi)
 - Overall system utilization $(U) = \sum ui$

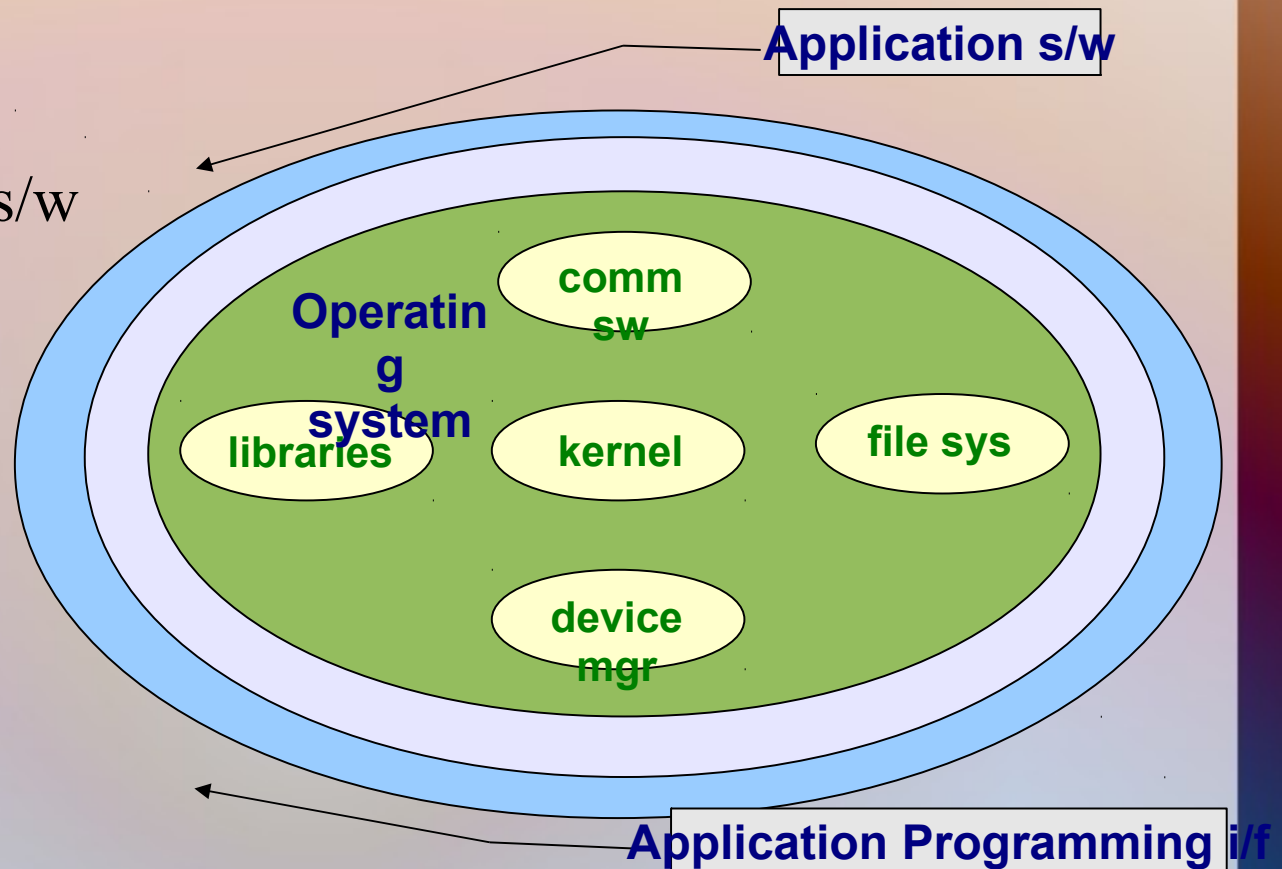
| Utilization (%) | Zone Type | Typical Application |
|-----------------|---|---------------------|
| 0-25 | significant excess processing power – CPU may be more powerful than necessary | various |
| 26-50 | very safe | various |
| 51-68 | safe | various |
| 69 | theoretical limit | embedded systems |
| 70-82 | questionable | embedded systems |
| 83-99 | dangerous | embedded systems |
| 100+ | overload | stressed systems |

RTOS Vs GPOS

| # | Metric of Evaluation | General Purpose OS | Real Time OS |
|---|-------------------------|---|--|
| 1 | Determinism | Non Deterministic | Deterministic – Kernel functions should execute in a fixed amount of time |
| 2 | Load Independent Timing | Not Applicable – Response becomes sluggish as number of tasks increase | Remains Constant – Generally Priority based execution. Time Slices are available only among tasks that hold the same priority |
| 3 | Task Level Scheduling | Generally Round Robin Scheduling, Sometimes Priority based scheduling – Efforts are made to ensure that all tasks get a chance to execute | Priority based Preemptive Scheduling – Ensure that the Highest Priority task is always executed, even if it is the most frequent |
| 4 | Interrupt Management | Nesting may be disabled. Interrupt latency, response and recovery are not performance metrics | Nesting is always enabled. Interrupt latency, response and recovery are very important performance metrics |

RTOS Architecture


- An RTOS consists of
 - Kernel (micro-kernel)
 - Device Manager
 - Networking protocol s/w
 - Libraries
 - File sys (optional)



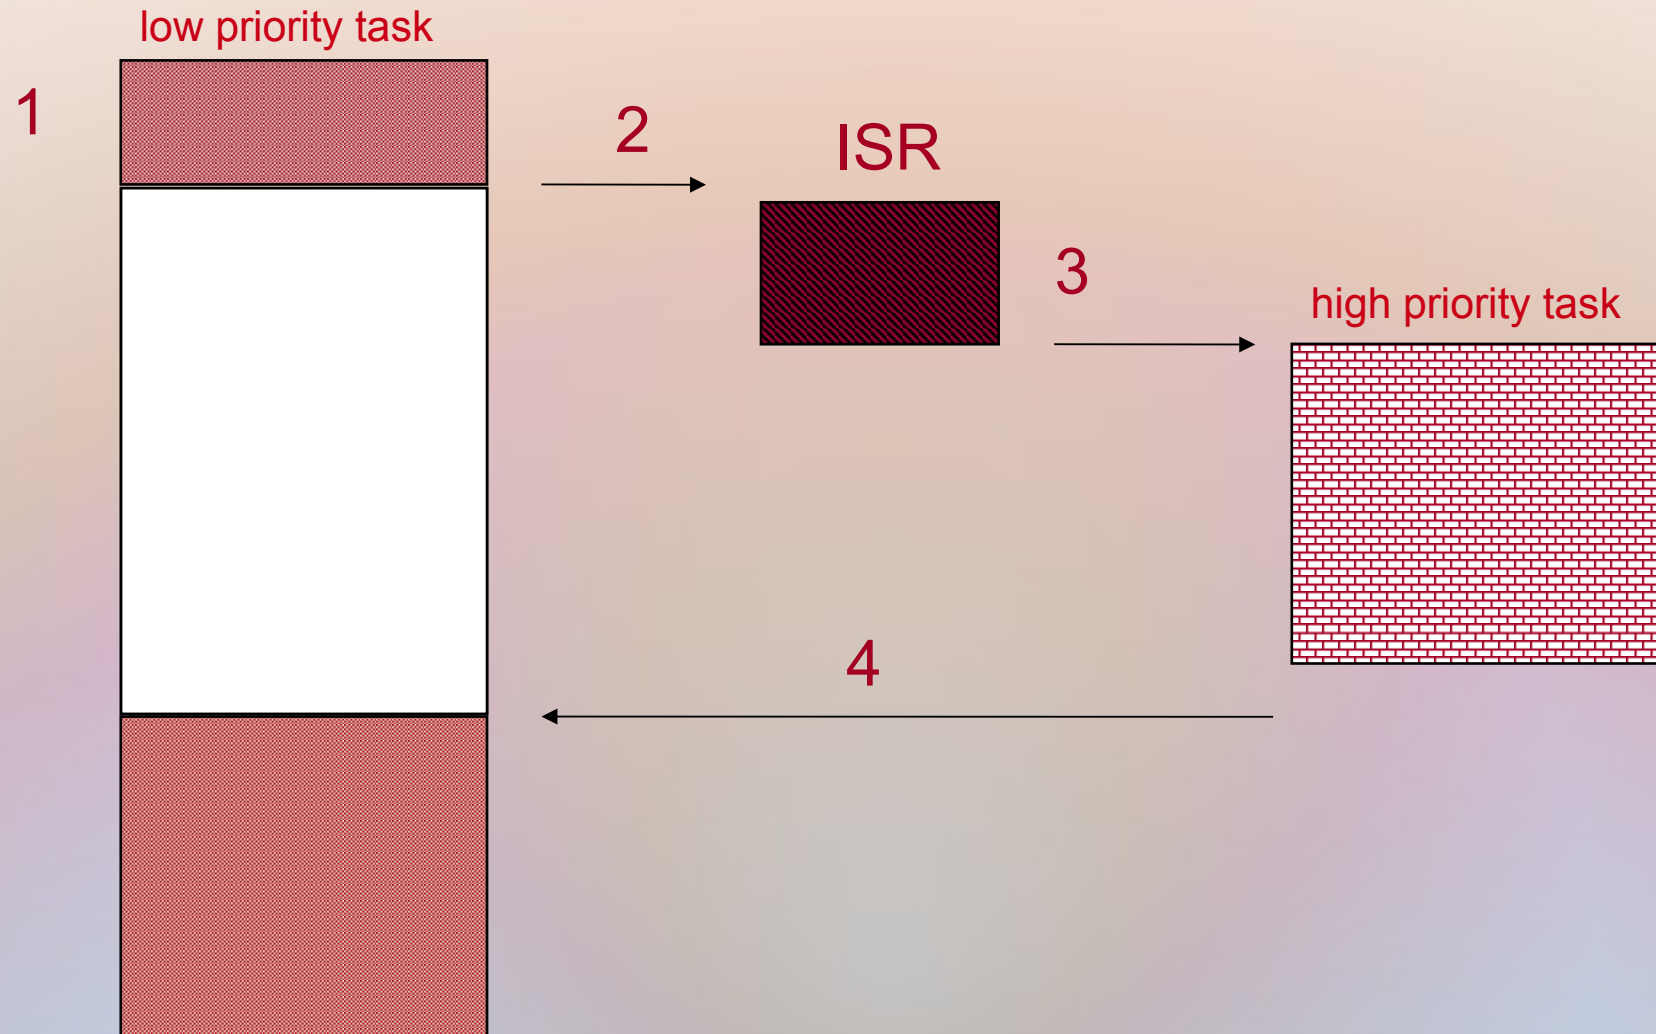
RTOS Services

- Multi-tasking, Task & Thread Management
 - Creation and scheduling of tasks.
 - Priority based pre-emptive scheduling
- Inter Task Synchronization and Inter Task Communications through mutual exclusions, signals, messages, shared memory, etc
- Usually no Memory Management
- Timer Services such as periodic and aperiodic interrupts

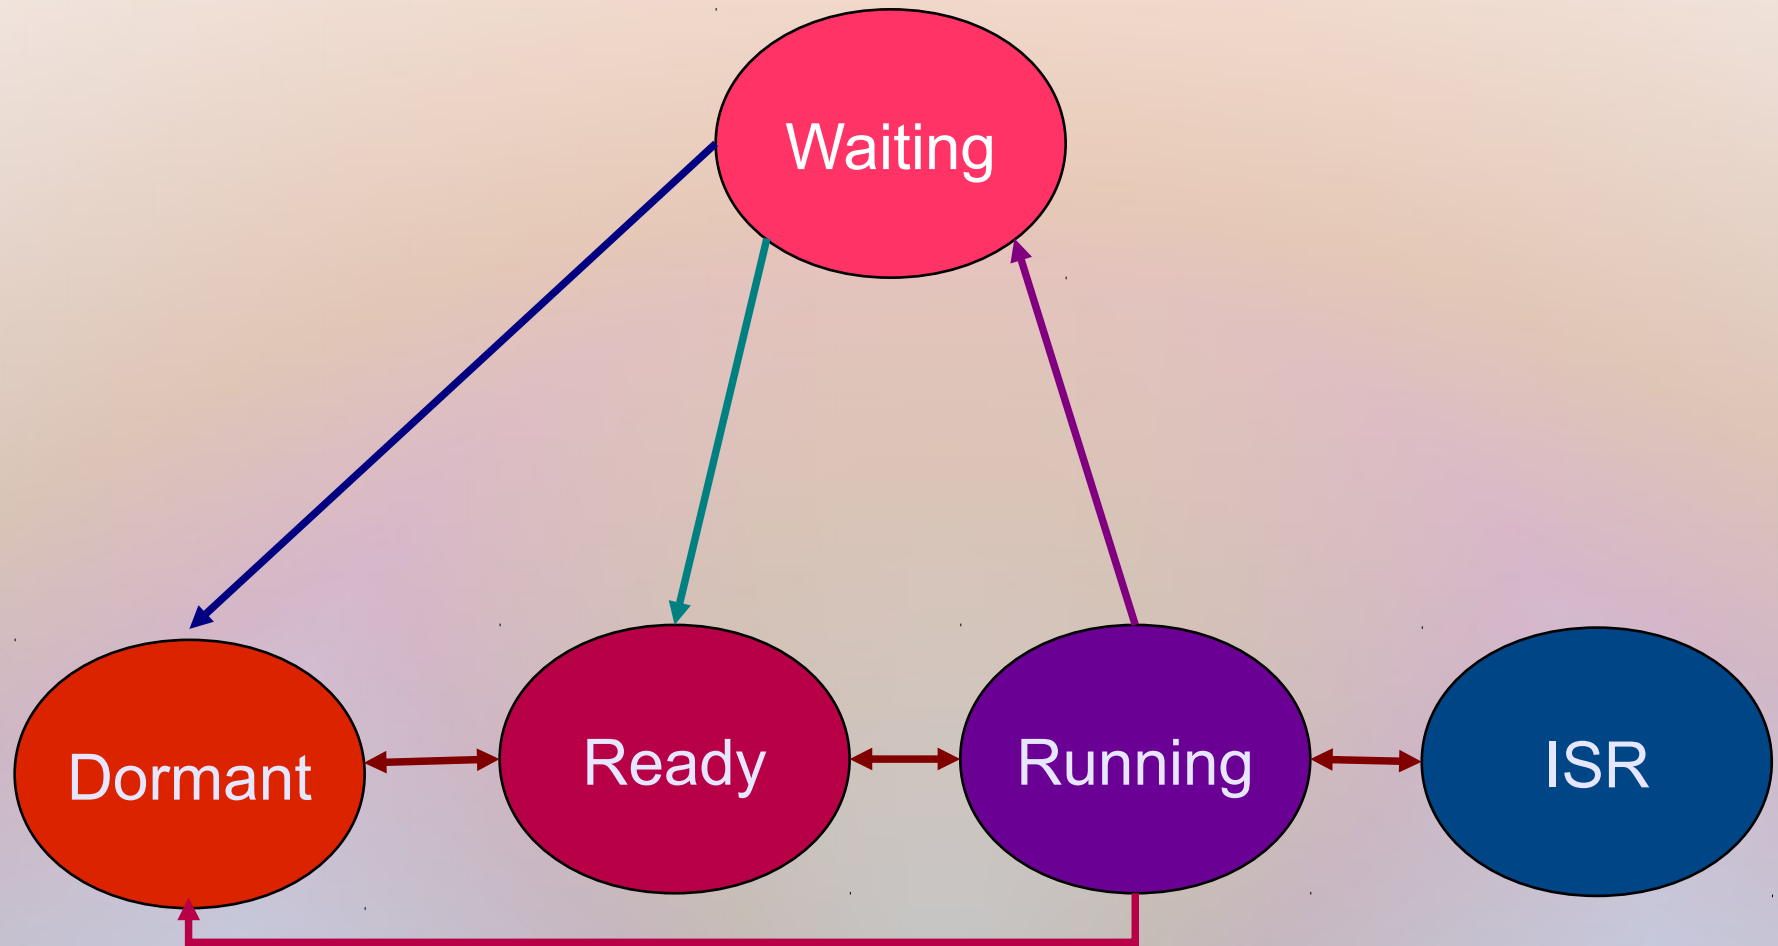
Kernel Types – Non Preemptive Kernel



Kernel Types – Preemptive Kernel



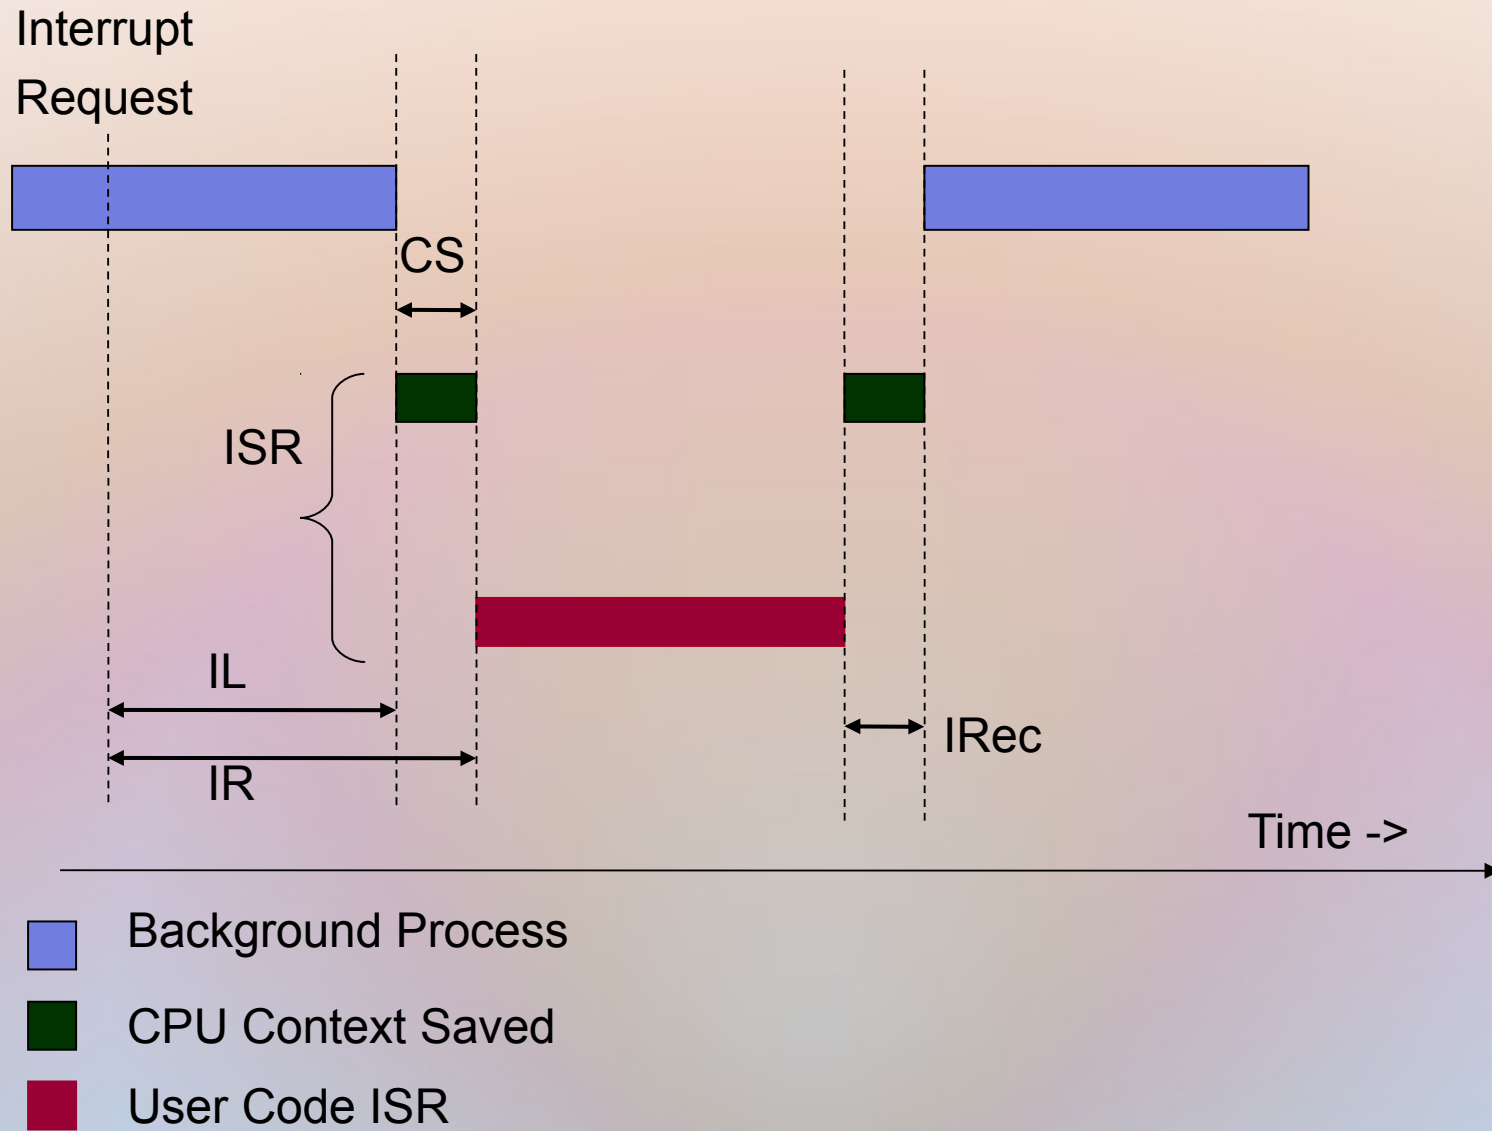
Task States



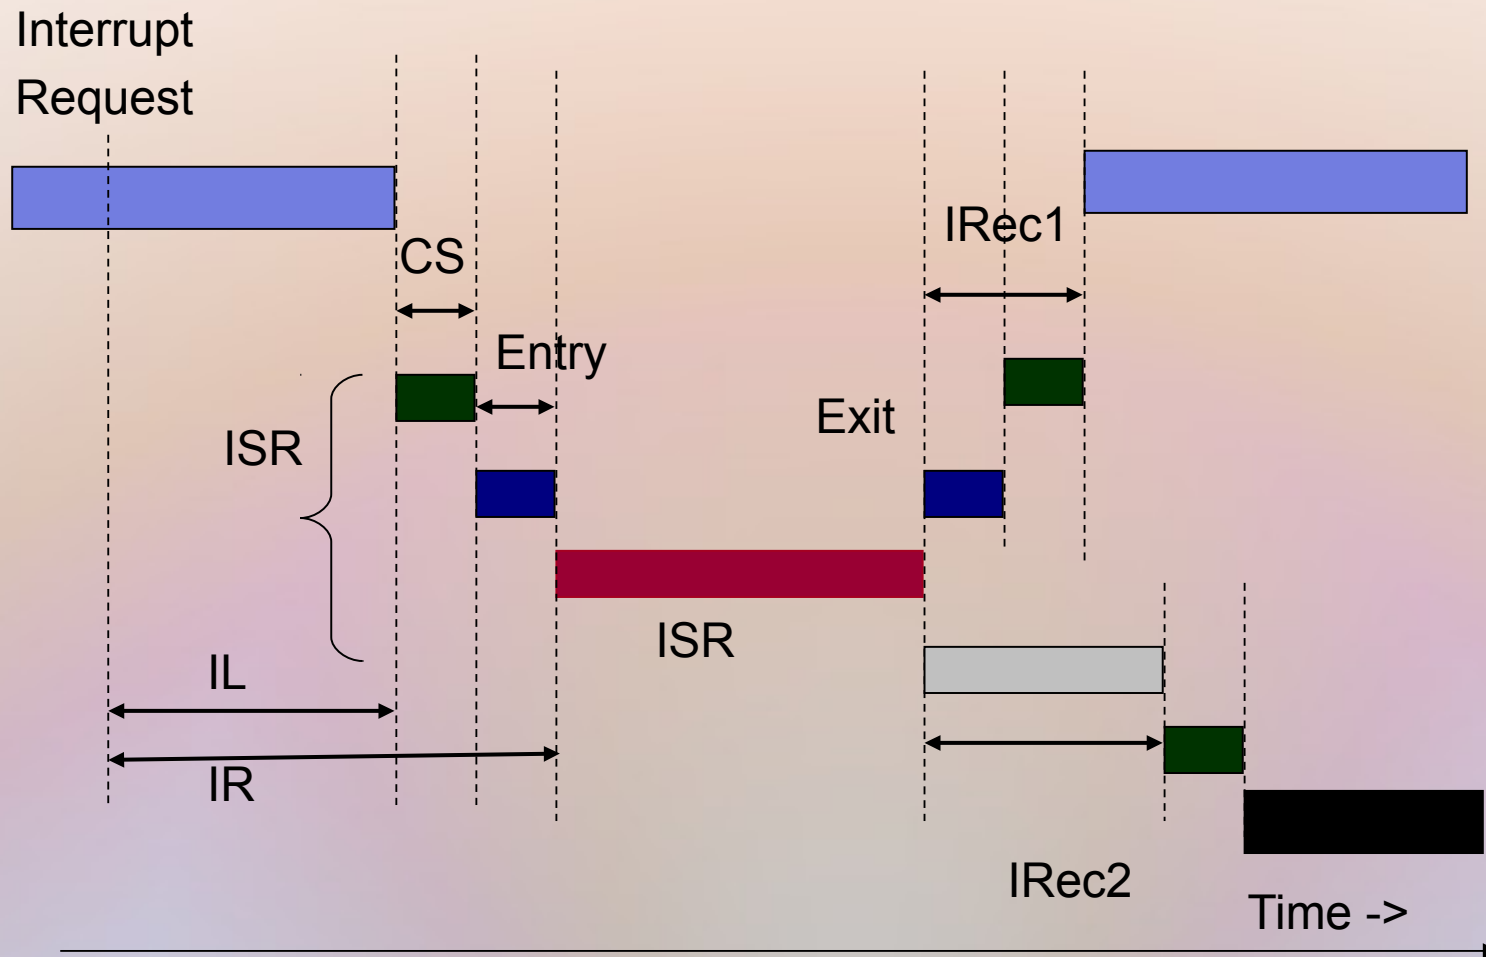
Interrupts

- When an interrupt occurs,
 - CPU saves its context on the stack, Jumps to the Interrupt Servicing Routine (ISR), Executes ISR and Returns
 - **Interrupt Latency**
 - max time interrupts are disabled + time to begin servicing the interrupt
 - **Interrupt Response Time**
 - Interrupt Latency + time to start execution of 1st instruction in ISR
 - **Interrupt Recovery Time**
 - time for CPU to return to interrupted code / highest priority task

Interrupts in Non Preemptive Kernels



Interrupts in a preemptive kernel



- Background Process
- CPU Context Saved
- User Code ISR

Preemptive Kernel


Priority Inversion, Inheritance and Ceiling

Task HP

Task MP


Task LP

 Shared Resource



Overview of available RTOS's

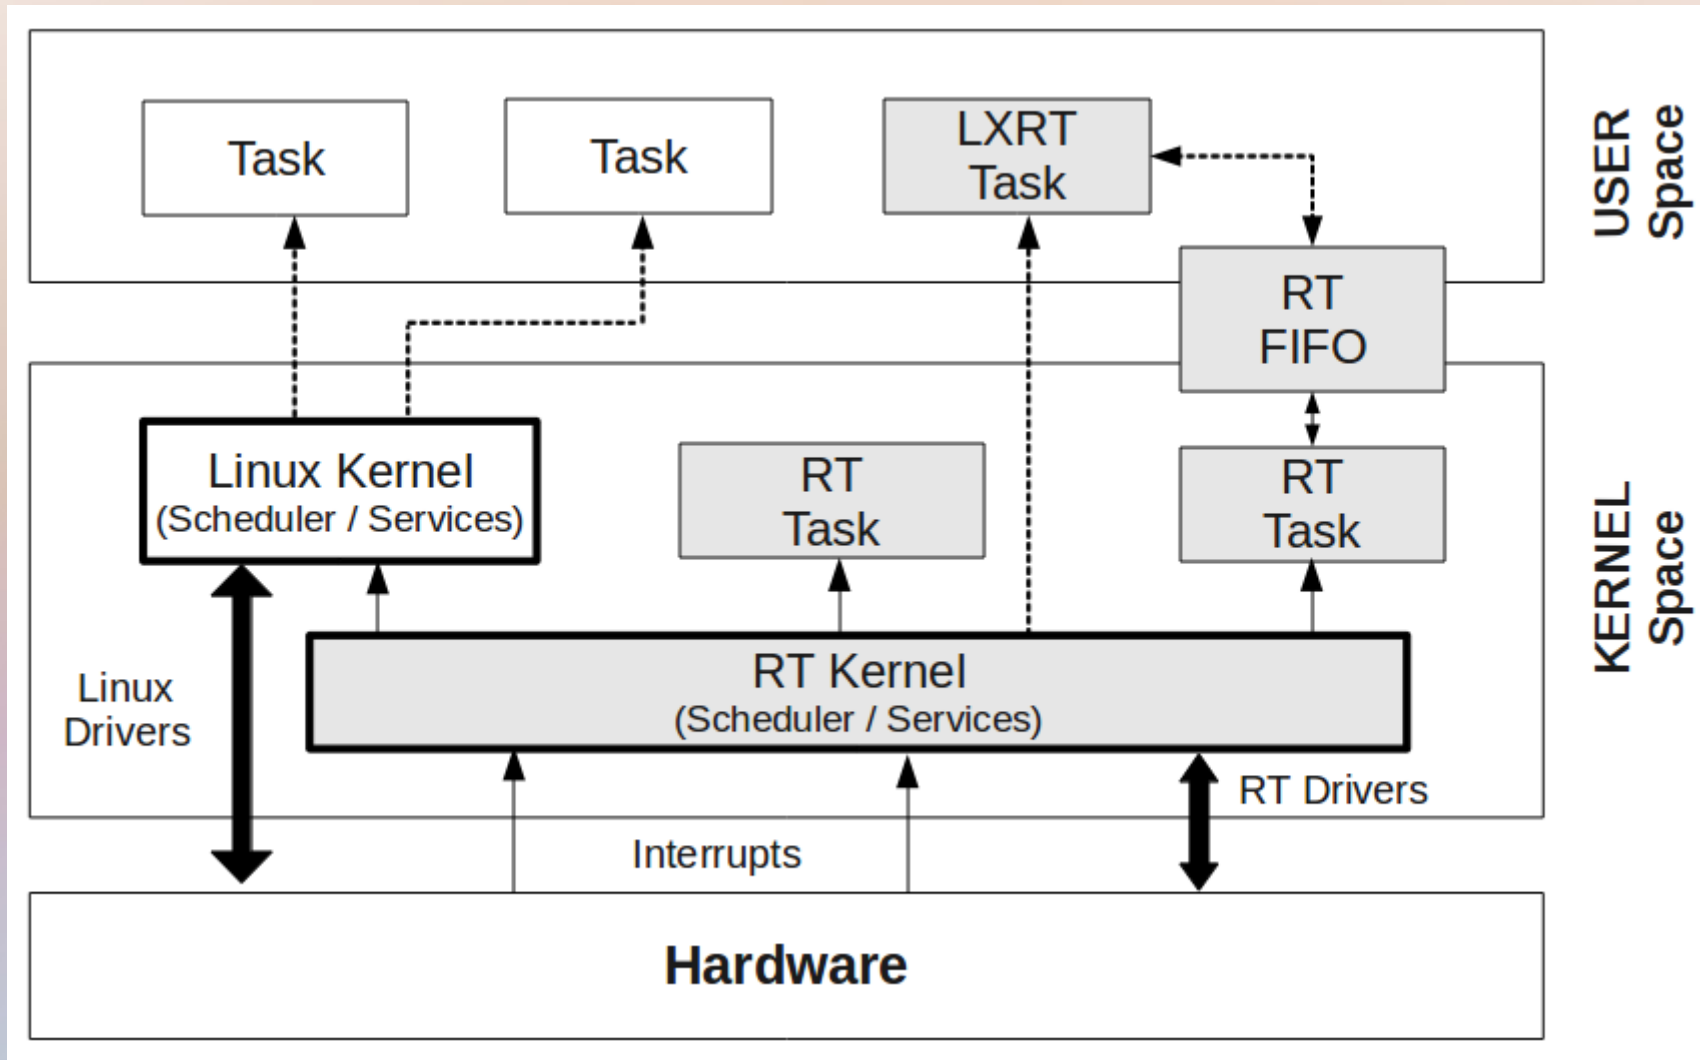
- Three categories of real-time operating systems:
 - Small, proprietary kernels. e.g. VRTX32, pSOS, VxWorks, Windows CE, MicroC-OS/III*
 - Real-time Linux extensions: RT-Linux, Xenomai, RTAI
 - Research kernels: MARS, ARTS, Spring, Polis, MicroC-OS/II



Variants of Real-Time Linux

- Two primary variants of hard real-time Linux available are RTLinux and RTAI.
 - **Real Time Linux (RTLinux)** was developed at the New Mexico Institute of Technology
 - www.rtlinux.org
 - **Real-Time Application Interface (RTAI)** was developed by programmers at the Department of Aerospace Engineering, Milano
 - www.aero.polimi.it/~rtai

Real Time Linux



The Real-Time Linux

- In a real-time Linux system
 - Real-time linux scheduler treats Linux kernel as the idle task
 - Linux only executes when there are no real time tasks to run, and when the real time kernel is inactive.
 - Linux is not permitted to disable hardware interrupts, I.e cannot add latency to the interrupt response time
 - When Linux code tries to disable interrupts, the real time system intercepts the request, records it. Doesn't actually disable ITRs
 - When an interrupt occurs, the real time kernel intercepts the interrupt and dispatches
 - First real time handlers are invoked then Linux handlers

The Real-Time Linux

- Real time kernel never waits for the Linux side to release any resources.
 - Communication links that are used to transfer data between real time tasks and Linux processes are non-blocking on the real time side.
- Linux operating system does as much as is practicable
 - System and Device initialization
 - Blocking dynamic resource allocation
 - Thread of execution that can be blocked
 - Loadable module mechanism to install components of the Real-Time system.

The Real-Time Linux

- The Real Time kernel, all its component parts, and the real time applications run in Linux kernel address space as kernel modules.
 - Advantage: Aids in modularity
 - Disadvantage: A bug in a real time task can crash the whole system.
- Real-time Linux decouples the services of the real time kernel from the services of the general purpose Linux kernel.
 - Real-time kernel can be kept small and simple, each service can be optimized independently

Real-Time Application Interface (RTAI)

- RTAI versions over 3.0 use an Adeos kernel patch, comprising an Interrupt Pipeline, where different Operating System Domains register interrupt handlers.
- Adeos (Adaptive Domain Environment for Operating Systems) is a nanokernel hardware abstraction layer (HAL) that operates between computer hardware and the operating system that runs on it.
 - Provides environment for sharing hardware resources among multiple OS, or multiple instances of a single OS

Real-Time Application Interface (RTAI)

- RTAI provides deterministic and preemptive performance in addition to allowing the use of all standard Linux drivers, applications and functions.
- RTAI's features:
 - Traditional RTOS IPCs including: Semaphores, mailboxes, FIFOs, shared memory, and RPCs
 - POSIX 1003.1c (Pthreads, mutexes and condition variables) & POSIX 1003.1b (Pqueues only) compatibility

Real-Time Application Interface (RTAI)

- RTAI's features:
 - /proc interface – which provides information on the real-time tasks, modules, services and processes extending the standard Linux /proc file-system support.
 - LXRT – which allows the use of the RTAI system calls from within standard user space
 - UniProcessor, Multi-UniProcessor and Symmetric Multi-processor (SMP) support
 - One-shot and periodic schedulers

RTAI Performance

- RTAI's performance is very competitive with the best commercial Real Time Operating Systems (such as VxWorks, QNX etc)
- Offers:
 - Context switch time: 4 uSec
 - Interrupt response: 20 uSec
 - 100 KHz periodic tasks
 - 30 KHz one-shot task rate

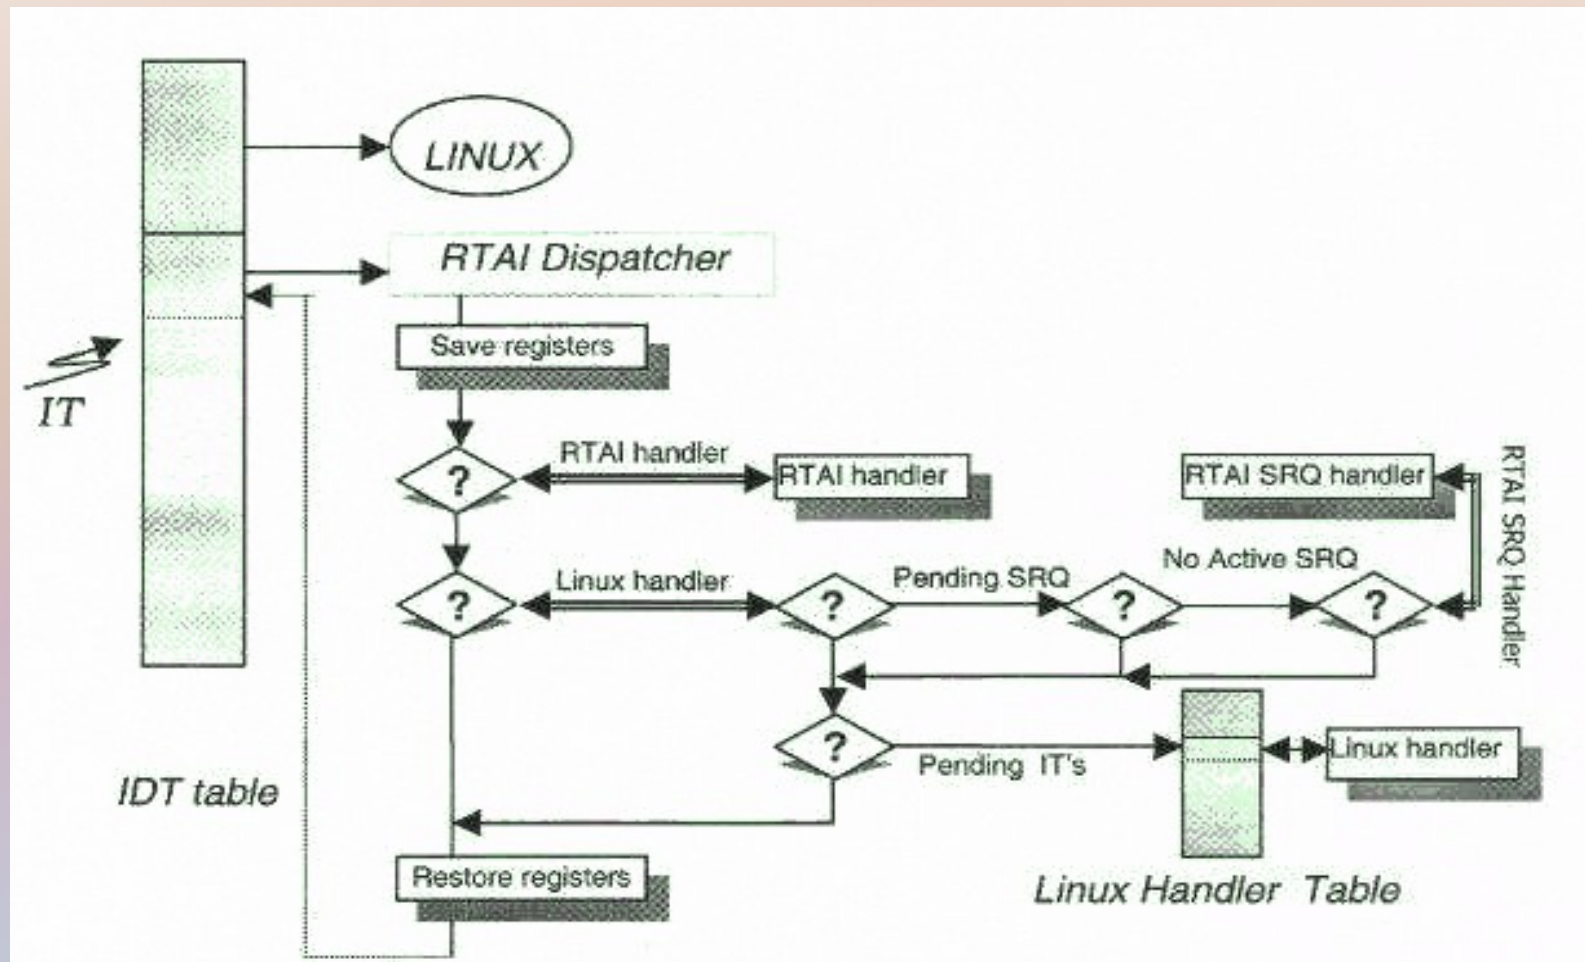
RTAI Components

- The three main RTAI components are:
 - The Interrupt Dispatcher
 - The Scheduler
 - The FIFO

RTAI Interrupt Dispatcher

- The dispatcher is a handler called via the IDT when interrupts occur.
 - Activates the right handler depending on the owner of the driver: RTAI, Linux or both (RTAI first then Linux).
 - The addresses of the Linux handlers are saved before changing the IDT in the HAL.
 - Takes into account the “SRQ(service request) handlers”.
 - SRQ is a handler used for a system call implementation. SRQ can be either RTAI handler (e.g.: fifo) or User handler.
 - RTAI manages up to 31 SRQ's. SRQ handler addresses are stored in IDT.

RTAI Interrupt Dispatcher



RTAI Scheduler

- The scheduler is in charge of distributing the CPU to different tasks present in the system (including Linux).
 - The RTAI distribution includes three different priority based, pre-emptive real time schedulers:
 - Uni-Processor (UP) scheduler;
 - Multi Uni-Processor (MUP) scheduler;
 - Symmetric Multi-Processor (SMP) scheduler
 - During the installation process, a scheduler is determined based on the hardware configuration of the target machine.

RTAI Scheduler

- UP scheduler - For uni-processor platforms.
 - Timer supports either one-shot or periodic scheduling but not both simultaneously.
- SMP scheduler - For multi-processor machines.
 - Timer supports either one-shot or periodic scheduling but not both simultaneously.
 - Tasks can run symmetrically on any or a cluster of CPUs, or be bound to a single CPU.
 - By default all tasks are defined to run on any of the CPUs and are automatically moved between CPUs as the system's processing and load requirements change.

RTAI Scheduler

- Multi-Uniprocessor scheduler - For multiprocessor platforms only and supports both one-shot and periodic scheduling simultaneously.
 - The main advantage is the ability to be able to use mixed timers simultaneously, i.e. periodic and one-shot timers.
 - Periodic timers can be based on different periods.
 - Like the SMP schedulers, the MUP can use inter-CPU services related to semaphores, messages and mailboxes.

Scheduler Functionality

- The RTAI multitasking scheduler, uses interrupt-driven, priority-based task scheduling.
- The scheduler elects the first highest priority task in a READY state.
 - Priority 0 is the highest priority and 0x3fffFfff the lowest for rtai tasks
 - Linux is given priority 0x7fffFfff.
 - RTAI takes care of Linux activation. When timer handler is running and the next Linux time period is reached a Linux pending interrupt is raised and will be served by the dispatcher.

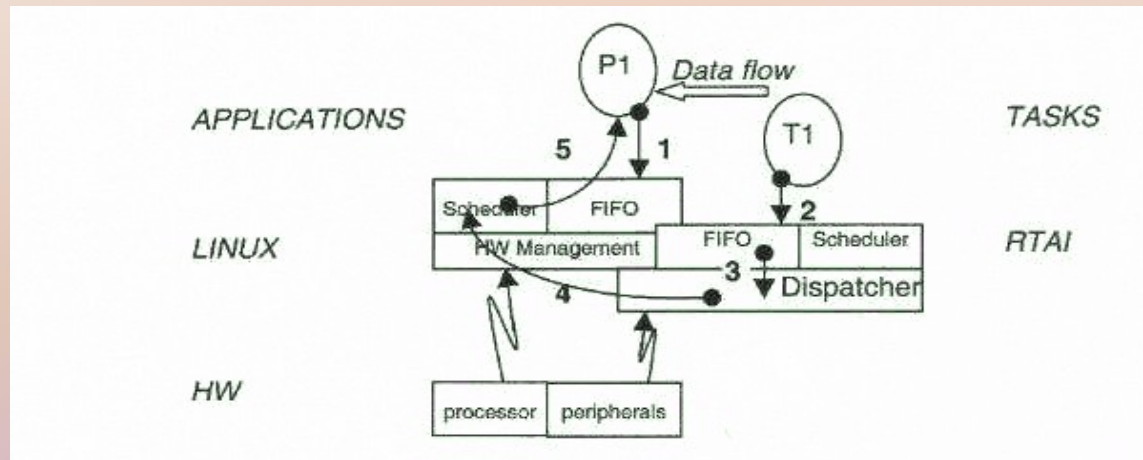
Scheduler Functionality

- All RTAI schedulers incorporate standard RTOS scheduling services like resume, yield, suspend, make periodic, wait until etc.
- The scheduler looks for ready tasks or tasks with expired period.
 - A task may be in the following states : READY, SUSPENDED, DELAYED, SEMAPHORE, SEND, RECEIVE, RPC, RETURN, RUNNING, DELETED
 - On Timer event, the timer handler changes the task state.

RTAI FIFOs

- Fifos allow Linux processes and RTAI tasks to exchange byte-oriented data streams.
 - A fifo is a one way channel therefore a duplex communication needs 2 fifos.
 - From a Linux process point of view the mechanism allows either blocking or non-blocking IO depending on user program.
 - From a task point of view fifo allows non-blocking or asynchronous IO.
 - To perform synchronous reads (and to avoid polling) RTAI makes it possible to attach a user real time handler to a fifo.

RTAI FIFOs



- blocking read.
- `rtf_put` call.
- SRQ initialization.
- wake-up action in queue.
- wake-up action de-queued at next Linux scheduling

RTAI LXRT Module

- LXRT is a module that allows to use all the services made available by RTAI and its schedulers in user space, both for soft and hard real time.